

# Deep Learning

QCon London, March 9th

# Today's Overview

# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch

# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch
- ▶ 10:30 Coffee Break

# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch
- ▶ 10:30 Coffee Break
- ▶ 10:45 Neural Networks
  - ▶ Implement a multi-layer perceptron
  - ▶ Recurrent neural networks

# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch
- ▶ 10:30 Coffee Break
- ▶ 10:45 Neural Networks
  - ▶ Implement a multi-layer perceptron
  - ▶ Recurrent neural networks
- ▶ 12:00 Lunch

# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch
- ▶ 10:30 Coffee Break
- ▶ 10:45 Neural Networks
  - ▶ Implement a multi-layer perceptron
  - ▶ Recurrent neural networks
- ▶ 12:00 Lunch
- ▶ 13:00 Word2Vec

# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch
- ▶ 10:30 Coffee Break
- ▶ 10:45 Neural Networks
  - ▶ Implement a multi-layer perceptron
  - ▶ Recurrent neural networks
- ▶ 12:00 Lunch
- ▶ 13:00 Word2Vec
- ▶ 14:00 Chatbot



# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch
- ▶ 10:30 Coffee Break
- ▶ 10:45 Neural Networks
  - ▶ Implement a multi-layer perceptron
  - ▶ Recurrent neural networks
- ▶ 12:00 Lunch
- ▶ 13:00 Word2Vec
- ▶ 14:00 Chatbot
- ▶ 14:30 Coffee Break

# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch
- ▶ 10:30 Coffee Break
- ▶ 10:45 Neural Networks
  - ▶ Implement a multi-layer perceptron
  - ▶ Recurrent neural networks
- ▶ 12:00 Lunch
- ▶ 13:00 Word2Vec
- ▶ 14:00 Chatbot
- ▶ 14:30 Coffee Break
- ▶ 14:45 Chatbot (continued)

# Today's Overview

- ▶ 9:00 Introduction to Deep Learning
  - ▶ Short recap on machine learning
  - ▶ Build and train a perceptron in numpy
  - ▶ Gradient Descent
  - ▶ Move the code to the GPU using PyTorch
- ▶ 10:30 Coffee Break
- ▶ 10:45 Neural Networks
  - ▶ Implement a multi-layer perceptron
  - ▶ Recurrent neural networks
- ▶ 12:00 Lunch
- ▶ 13:00 Word2Vec
- ▶ 14:00 Chatbot
- ▶ 14:30 Coffee Break
- ▶ 14:45 Chatbot (continued)
- ▶ 16:00 End of workshop

# Why Deep Learning?

# Why Deep Learning?

- ▶ It can save a lot of time of booooooring feature engineering and writing exceptions

# Why Deep Learning?

- ▶ It can save a lot of time of booooooring feature engineering and writing exceptions
- ▶ It can write computational rules nobody of us would ever think of

# Why Deep Learning?

- ▶ It can save a lot of time of booooooring feature engineering and writing exceptions
- ▶ It can write computational rules nobody of us would ever think of
- ▶ It can often write better computational rules than we could


# Why Deep Learning?

- ▶ It can save a lot of time of booooooring feature engineering and writing exceptions
- ▶ It can write computational rules nobody of us would ever think of
- ▶ It can often write better computational rules than we could
- ▶ It can check more cases than we could




# Perceptron

- ▶ The perceptron contains a weight ( $\vec{w}$ ) for each input and a threshold ( $\Theta$ )




# Perceptron

- ▶ The perceptron contains a weight ( $\vec{w}$ ) for each input and a threshold ( $\Theta$ )



# Perceptron

- ▶ The perceptron contains a weight ( $\vec{w}$ ) for each input and a threshold ( $\Theta$ )
- ▶ Its output can be calculated as  $\sum_{i=0}^k w_i X_i \geq \Theta$



# Perceptron Training

# Perceptron Training

- ▶ Start with random weights

# Perceptron Training

- ▶ Start with random weights
- ▶ Calculate the perceptron's estimate of the solution with:  $\hat{y} = \left( \sum_i w_i X_i \geq 0 \right)$

# Perceptron Training

- ▶ Start with random weights
- ▶ Calculate the perceptron's estimate of the solution with:  $\hat{y} = \left( \sum_i w_i X_i \geq 0 \right)$
- ▶ Determine the update for the weights using the error and the learning rate ( $\alpha$ ):  
 $\partial w_i = \alpha(y - \hat{y})X_i$

# Perceptron Training

- ▶ Start with random weights
- ▶ Calculate the perceptron's estimate of the solution with:  $\hat{y} = \left( \sum_i w_i X_i \geq 0 \right)$
- ▶ Determine the update for the weights using the error and the learning rate ( $\alpha$ ):  
 $\partial w_i = \alpha(y - \hat{y})X_i$
- ▶ Calculate new weights as:  $w_i \leftarrow w_i + \partial w_i$



# Perceptron Training

- ▶ Start with random weights
- ▶ Calculate the perceptron's estimate of the solution with:  $\hat{y} = \left( \sum_i w_i X_i \geq 0 \right)$
- ▶ Determine the update for the weights using the error and the learning rate ( $\alpha$ ):  
 $\partial w_i = \alpha(y - \hat{y})X_i$
- ▶ Calculate new weights as:  $w_i \leftarrow w_i + \partial w_i$
- ▶ Restart the process until the solution is found

# Perceptron Training

- ▶ Start with random weights
- ▶ Calculate the perceptron's estimate of the solution with:  $\hat{y} = \left( \sum_i w_i X_i \geq 0 \right)$
- ▶ Determine the update for the weights using the error and the learning rate ( $\alpha$ ):  
 $\partial w_i = \alpha(y - \hat{y})X_i$
- ▶ Calculate new weights as:  $w_i \leftarrow w_i + \partial w_i$
- ▶ Restart the process until the solution is found
- ▶ Let us actually implement this in the first notebook (Perceptron)

# Conclusions


# Conclusions

- ▶ Congratulations on training your first perceptron! You just taught a computer to learn!


# Conclusions

- ▶ Congratulations on training your first perceptron! You just taught a computer to learn!
- ▶ Why didn't we wait for the final weights, but stop after one run over the dataset?

# Linear Separability


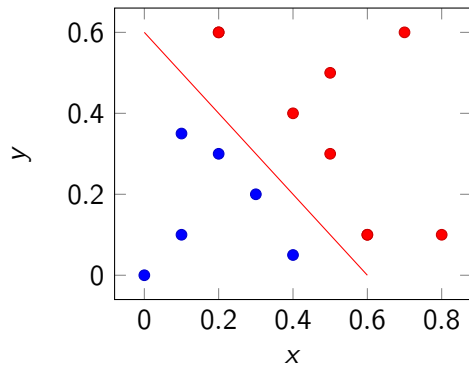


## Linear Separability



Data are linearly separable -> Can be divided by a plane


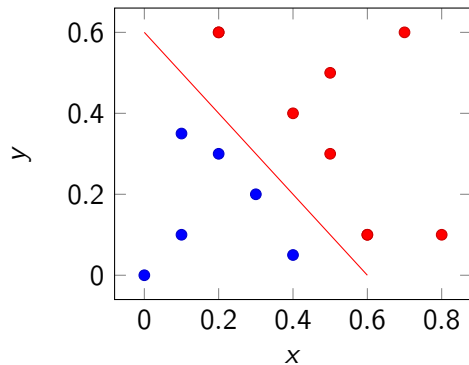
## Linear Separability



Data are linearly separable -> Can be divided by a plane


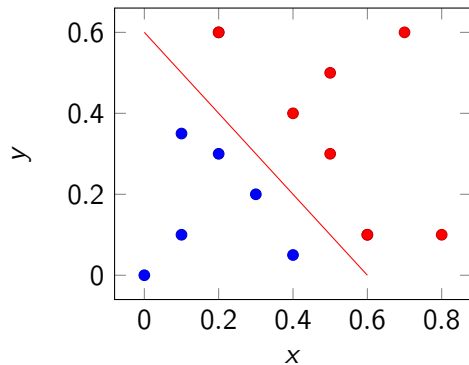


## Linear Separability




Data are linearly separable -> Can be divided by a plane

## Linear Separability




Data are linearly separable -> Can be divided by a plane

## Linear Separability




Data are linearly separable -> Can be divided by a plane




Data is not linearly separable -> Can not be divided by a plane

## Linear Separability



Data are linearly separable -> Can be divided by a plane



Data is not linearly separable -> Can not be divided by a plane

The algorithm only converges, if the problem is linearly separable.

# Gradient Descent

# Gradient Descent

- ▶ What we calculated is an activation

$$a = \sum_i w_i X_i$$

# Gradient Descent


- ▶ What we calculated is an activation

$$a = \sum_i w_i X_i$$

- ▶ Let us ignore the threshold and try to get the activation as close to the target value as possible.

## Gradient Descent


- ▶ What we calculated is an activation  
$$a = \sum_i w_i X_i$$
- ▶ Let us ignore the threshold and try to get the activation as close to the target value as possible.
- ▶ This brings us back to regression with an error:  $E(w) = \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2$






# Gradient Descent

- ▶ What we calculated is an activation  
$$a = \sum_i w_i X_i$$
- ▶ Let us ignore the threshold and try to get the activation as close to the target value as possible.
- ▶ This brings us back to regression with an error:  $E(w) = \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2$




# Gradient Descent

- ▶ What we calculated is an activation  
$$a = \sum_i w_i X_i$$
- ▶ Let us ignore the threshold and try to get the activation as close to the target value as possible.
- ▶ This brings us back to regression with an error:  $E(w) = \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2$
- ▶ If we want to decrease  $E(w)$  by changing  $w$  we need to calculate the gradient of  $w$



# Gradient Descent

- ▶ What we calculated is an activation  
$$a = \sum_i w_i X_i$$
- ▶ Let us ignore the threshold and try to get the activation as close to the target value as possible.
- ▶ This brings us back to regression with an error:  $E(w) = \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2$
- ▶ If we want to decrease  $E(w)$  by changing  $w$  we need to calculate the gradient of  $w$
- ▶ Every deep learning framework can calculate those weights using the chain rule and backpropagation.



# Optimizing Weights

- ▶ There are multiple ways to use gradient descent.

# Optimizing Weights

- ▶ There are multiple ways to use gradient descent.
- ▶ "Simple" Gradient Descent

# Optimizing Weights

- ▶ There are multiple ways to use gradient descent.
- ▶ "Simple" Gradient Descent
- ▶ Advanced Methods

# Optimizing Weights

- ▶ There are multiple ways to use gradient descent.
- ▶ "Simple" Gradient Descent
- ▶ Advanced Methods
  - ▶ Momentum

# Optimizing Weights

- ▶ There are multiple ways to use gradient descent.
- ▶ "Simple" Gradient Descent
- ▶ Advanced Methods
  - ▶ Momentum
  - ▶ Higher Order Derivatives



# Optimizing Weights

- ▶ There are multiple ways to use gradient descent.
- ▶ "Simple" Gradient Descent
- ▶ Advanced Methods
  - ▶ Momentum
  - ▶ Higher Order Derivatives
  - ▶ Randomised Optimisation

# Optimizing Weights

- ▶ There are multiple ways to use gradient descent.
- ▶ "Simple" Gradient Descent
- ▶ Advanced Methods
  - ▶ Momentum
  - ▶ Higher Order Derivatives
  - ▶ Randomised Optimisation
  - ▶ Penalty for Complexity


Sigmoid  $\rightarrow$  Differentiable Threshold

## Sigmoid $\rightarrow$ Differentiable Threshold

- ▶ Even though the perceptron is great for boolean functions it is not differentiable and has no continuous error function, consequently the gradient cannot be calculated


## Sigmoid $\rightarrow$ Differentiable Threshold

- ▶ Even though the perceptron is great for boolean functions it is not differentiable and has no continuous error function, consequently the gradient cannot be calculated
- ▶  $\sigma(a) = \frac{1}{1+e^{-a}}$



## Sigmoid $\rightarrow$ Differentiable Threshold

- ▶ Even though the perceptron is great for boolean functions it is not differentiable and has no continuous error function, consequently the gradient cannot be calculated
- ▶  $\sigma(a) = \frac{1}{1+e^{-a}}$
- ▶ The sigmoid allows us to convert the strict classification into a regression over the probability for each point to belong to a certain class



# Softmax

# Softmax

- ▶ To generalize to n-class classification problems we have to extend the function, that calculates the probability



# Softmax

- ▶ To generalize to n-class classification problems we have to extend the function, that calculates the probability
- ▶ Softmax:  $p_i = \frac{e^{o_i}}{\sum_i e^{o_i}}$  probability for being in class i, given the output  $o_i$

# Softmax

- ▶ To generalize to n-class classification problems we have to extend the function, that calculates the probability
- ▶ Softmax:  $p_i = \frac{e^{o_i}}{\sum_i e^{o_i}}$  probability for being in class i, given the output  $o_i$
- ▶ With this we can define the general loss

# Softmax

- ▶ To generalize to n-class classification problems we have to extend the function, that calculates the probability
- ▶ Softmax:  $p_i = \frac{e^{o_i}}{\sum_i e^{o_i}}$  probability for being in class  $i$ , given the output  $o_i$
- ▶ With this we can define the general loss
- ▶ Cross-Entropy  $E = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^m y_{i,k} \ln p_{i,k}$   
 $y_{i,k}$  1, if  $i = k$  else 0  
 $p_{i,k}$  probability for point  $i$  belonging to class  $k$

# Neural Network


# Neural Network

- ▶ Similar to the perceptron we have a lot of inputs and output




# Neural Network

- ▶ Similar to the perceptron we have a lot of inputs and output
- ▶ However, instead of going from start to end directly we now add hidden units




# Neural Network

- ▶ Similar to the perceptron we have a lot of inputs and output
- ▶ However, instead of going from start to end directly we now add hidden units
- ▶ For each of the units we now apply the sigmoid function



# Neural Network

- ▶ Similar to the perceptron we have a lot of inputs and output
- ▶ However, instead of going from start to end directly we now add hidden units
- ▶ For each of the units we now apply the sigmoid function
- ▶ Since every unit is differentiable, the whole network is differentiable  
→ We can use backpropagation to minimize the error as we can calculate the gradients





# Restriction Bias

# Restriction Bias

- ▶ Representational power of the neural network (all hypothesis it can consider)

# Restriction Bias

- ▶ Representational power of the neural network (all hypothesis it can consider)
- ▶ Perceptron  $\rightarrow$  can only split the data using a planes

# Restriction Bias

- ▶ Representational power of the neural network (all hypothesis it can consider)
- ▶ Perceptron  $\rightarrow$  can only split the data using a planes
- ▶ Sigmoids  $\rightarrow$  non-linear function

# Restriction Bias

- ▶ Representational power of the neural network (all hypothesis it can consider)
- ▶ Perceptron  $\rightarrow$  can only split the data using a planes
- ▶ Sigmoids  $\rightarrow$  non-linear function
- ▶ What can we represent with a network?

# Restriction Bias

- ▶ Representational power of the neural network (all hypothesis it can consider)
- ▶ Perceptron  $\rightarrow$  can only split the data using a planes
- ▶ Sigmoids  $\rightarrow$  non-linear function
- ▶ What can we represent with a network?
  - ▶ Boolean Functions: Yes, because we have a network of threshold-like units

# Restriction Bias

- ▶ Representational power of the neural network (all hypothesis it can consider)
- ▶ Perceptron → can only split the data using a planes
- ▶ Sigmoids → non-linear function
- ▶ What can we represent with a network?
  - ▶ Boolean Functions: Yes, because we have a network of threshold-like units
  - ▶ Continuous Functions: Yes, with one (infinitely wide) layer

# Restriction Bias

- ▶ Representational power of the neural network (all hypothesis it can consider)
- ▶ Perceptron  $\rightarrow$  can only split the data using a planes
- ▶ Sigmoids  $\rightarrow$  non-linear function
- ▶ What can we represent with a network?
  - ▶ Boolean Functions: Yes, because we have a network of threshold-like units
  - ▶ Continuous Functions: Yes, with one (infinitely wide) layer
  - ▶ Arbitrary Functions: Yes, with two (infinitely wide) layers



# Restriction Bias

- ▶ Representational power of the neural network (all hypothesis it can consider)
- ▶ Perceptron → can only split the data using a planes
- ▶ Sigmoids → non-linear function
- ▶ What can we represent with a network?
  - ▶ Boolean Functions: Yes, because we have a network of threshold-like units
  - ▶ Continuous Functions: Yes, with one (infinitely wide) layer
  - ▶ Arbitrary Functions: Yes, with two (infinitely wide) layers

How do we make sure that we do not learn the arbitrary noise in the dataset?

# Regularization

# Regularization

- ▶ How can we limit the representational power of the neural network?

# Regularization

- ▶ How can we limit the representational power of the neural network?
- ▶ Limit the number of neurons

# Regularization

- ▶ How can we limit the representational power of the neural network?
- ▶ Limit the number of neurons
- ▶ Dropout


# Regularization

- ▶ How can we limit the representational power of the neural network?
- ▶ Limit the number of neurons
- ▶ Dropout
- ▶ Weight Decay

# Regularization

- ▶ How can we limit the representational power of the neural network?
- ▶ Limit the number of neurons
- ▶ Dropout
- ▶ Weight Decay
- ▶ All of these can be perfectly tested using a validation set


# Recurrent Neural Networks





# Recurrent Neural Networks


- Recurrent Neural Networks allow one to propagate states through time



# Recurrent Neural Networks


- ▶ Recurrent Neural Networks allow one to propagate states through time
- ▶ They calculate a hidden state  $\vec{h}_t$  which is kept for the next time step, as

$$h_t = \sigma_h(W_h x_t + U_h h_{t-1} + b_h)$$



# Recurrent Neural Networks

- ▶ Recurrent Neural Networks allow one to propagate states through time
- ▶ They calculate a hidden state  $\vec{h}_t$  which is kept for the next time step, as
$$\vec{h}_t = \sigma_h(W_h \vec{x}_t + U_h \vec{h}_{t-1} + b_h)$$
- ▶ From this hidden state the current output is calculated as
$$\vec{y}_t = \sigma_y(W_y \vec{h}_t + b_y)$$
- ▶ These networks cannot represent long-term dependencies as  $U_h$  is always multiplied by itself.



# LSTM

# LSTM

- ▶ LSTM (Long-Short-Term Memory) can store long term dependencies by adding a cell state and utilizing an input, forget and output gates.

# LSTM

- ▶ LSTM (Long-Short-Term Memory) can store long term dependencies by adding a cell state and utilizing an input, forget and output gates.
- ▶ The input gate decides how much of the input and hidden state are stored to the cell state.

# LSTM

- ▶ LSTM (Long-Short-Term Memory) can store long term dependencies by adding a cell state and utilizing an input, forget and output gates.
- ▶ The input gate decides how much of the input and hidden state are stored to the cell state.
- ▶ The forget gate decides how much of the cell state shall be deleted.

# LSTM

- ▶ LSTM (Long-Short-Term Memory) can store long term dependencies by adding a cell state and utilizing an input, forget and output gates.
- ▶ The input gate decides how much of the input and hidden state are stored to the cell state.
- ▶ The forget gate decides how much of the cell state shall be deleted.
- ▶ The output gate decides how much of the cell state is output to the hidden state.



# Word2Vec

# Word2Vec

- ▶ How can we turn text into numbers?

# Chatbot

# Chatbot

► Let's chat